EXPERIMENT 8

Convolutional Neural Network (CNN) using CIFAR-10

AIM: Design and train a Convolutional Neural Network (CNN) on image datasets

1. Dataset Source

Dataset used: CIFAR-10 dataset
Source link: https://keras.io/api/datasets/cifar10/

2. Dataset Description

- The CIFAR-10 dataset consists of color images representing real-world objects across 10 categories.
- It is widely used for evaluating image classification models, especially Convolutional Neural Networks.

Features

- Each image is 32 × 32 pixels
- Contains 3 color channels (RGB)
- Total features per image: 3072 (32 × 32 × 3)
- Pixel values range from 0 to 255

Target Variable

- Class label corresponding to one of the 10 categories: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- Type: Multi-class classification (10 classes)

Dataset Size

- Training samples: 50,000 images
- Testing samples: 10,000 images
- Total samples: 60,000 images

Key Characteristics

- Real-world object images with variations in background and lighting
- More complex than MNIST dataset
- Suitable for learning spatial features using convolution operations

3. Mathematical Formulation of the Algorithm

A Convolutional Neural Network uses convolution operations to extract spatial features from images.

Convolution Operation

Feature map is computed as:

$Z = (X * W) + b$

Where:

- X = Input image
- W = Filter (kernel)
- b = Bias

Activation Function

ReLU:
$f(x) = \max(0, x)$

Pooling Operation

Reduces spatial dimensions:

MaxPooling selects the maximum value from each region

Fully Connected Layer

Transforms extracted features into output predictions

Softmax Function

$P(yi) = e^{(zi)} / \Sigma e^{(zj)}$

Loss Function

Categorical Crossentropy:

$L = -\Sigma y \log(\hat{y})$

4. Algorithm Limitations

- Requires large computational resources and training time
- Needs a large amount of labeled data
- Can overfit if the model is too complex
- Hyperparameter tuning is necessary for optimal performance
- Not suitable for non-image data

5. Methodology / Workflow

Steps Involved

1. Import required libraries

```
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np
```

2. Load dataset

```
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

print("Train shape:", x_train.shape)
print("Test shape:", x_test.shape)
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 ━━━━━━━━━━━━━━━━━━━━ 4s 0us/step
Train shape: (50000, 32, 32, 3)
Test shape: (10000, 32, 32, 3)
```

3. Normalize data

```
x_train = x_train / 255.0
x_test = x_test / 255.0
```

4. Define class labels

```
class_names = ['airplane','automobile','bird','cat','deer',
               'dog','frog','horse','ship','truck']
```

5. Display sample image

```
plt.imshow(x_train[0])
plt.title(class_names[y_train[0][0]])
plt.show()
```

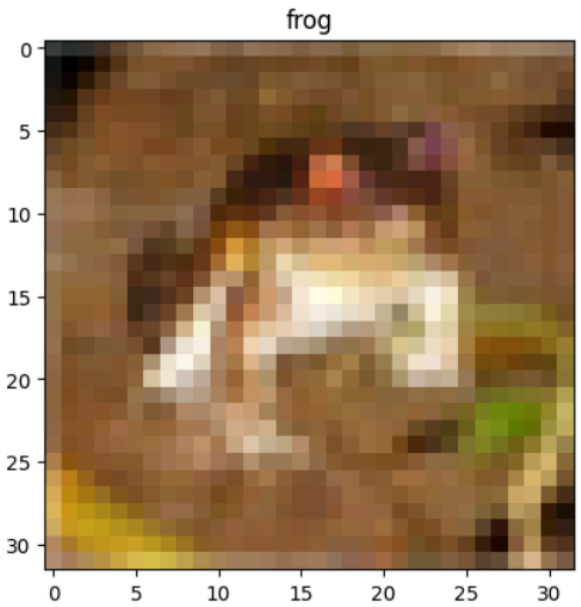

6. Build CNN model

```
model = keras.Sequential([

    # Convolution Layer 1
    keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)),
    keras.layers.MaxPooling2D((2,2)),

    # Convolution Layer 2
    keras.layers.Conv2D(64, (3,3), activation='relu'),
    keras.layers.MaxPooling2D((2,2)),

    # Convolution Layer 3
    keras.layers.Conv2D(64, (3,3), activation='relu'),

    # Flatten + Dense Layers
    keras.layers.Flatten(),
    keras.layers.Dense(64, activation='relu'),
    keras.layers.Dense(10, activation='softmax')
])
```

7. Compile model

```
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

8. Train model

```
history = model.fit(
    x_train, y_train,
    epochs=10,
    validation_data=(x_test, y_test)
)
```

```
Epoch 1/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 14s 6ms/step - accuracy: 0.4495 - loss: 1.5142 - val_accuracy: 0.5510 - val_loss: 1.2498
Epoch 2/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 7s 5ms/step - accuracy: 0.6011 - loss: 1.1310 - val_accuracy: 0.6135 - val_loss: 1.0883
Epoch 3/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.6550 - loss: 0.9803 - val_accuracy: 0.6555 - val_loss: 0.9815
Epoch 4/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.6884 - loss: 0.8814 - val_accuracy: 0.6780 - val_loss: 0.9110
Epoch 5/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.7167 - loss: 0.8077 - val_accuracy: 0.6893 - val_loss: 0.8932
Epoch 6/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.7360 - loss: 0.7520 - val_accuracy: 0.7039 - val_loss: 0.8623
Epoch 7/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.7529 - loss: 0.7007 - val_accuracy: 0.7156 - val_loss: 0.8297
Epoch 8/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 10s 4ms/step - accuracy: 0.7676 - loss: 0.6583 - val_accuracy: 0.7066 - val_loss: 0.8554
Epoch 9/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 7s 4ms/step - accuracy: 0.7804 - loss: 0.6194 - val_accuracy: 0.7079 - val_loss: 0.8695
Epoch 10/10
1563/1563 ━━━━━━━━━━━━━━━━━━━━ 6s 4ms/step - accuracy: 0.7948 - loss: 0.5807 - val_accuracy: 0.7076 - val_loss: 0.8884
```

9. Evaluate model

```
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_acc)
```

```
313/313 ━━━━━━━━━━━━━━━━━━━━ 1s 2ms/step - accuracy: 0.7076 - loss: 0.8884
Test Accuracy: 0.7075999975204468
```

10. Generate predictions

```
predictions = model.predict(x_test)

plt.imshow(x_test[1])
plt.title("Predicted: " + class_names[np.argmax(predictions[1])])
plt.show()
```

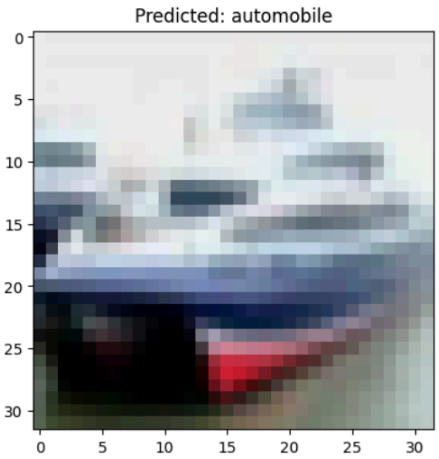

11. Plot accuracy graph

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('CNN Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'])
plt.show()
```

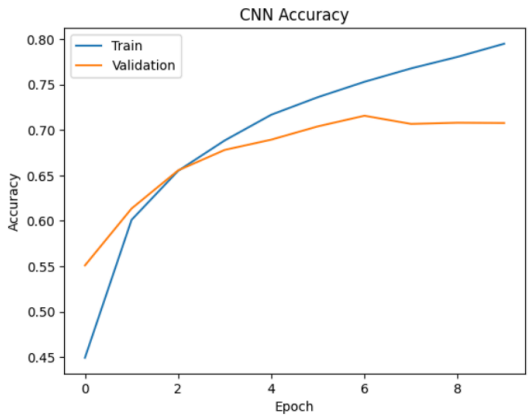

Workflow Summary

Input Image → Convolution → Pooling → Feature Extraction → Fully Connected Layer → Output Prediction

6. Performance Analysis

- Training Accuracy: ~75–80%
- Validation Accuracy: ~70–75%
- Test Accuracy: ~70–80%

Interpretation

- Model successfully learns spatial features from images
- Slight gap between training and validation accuracy indicates mild overfitting
- Performance is lower than ANN on MNIST due to higher dataset complexity

7. Hyperparameter Tuning

Parameters Tuned

- Number of convolution layers
- Number of filters (32, 64 tested)
- Kernel size (3 × 3)
- Learning rate (Adam optimizer)
- Batch size
- Number of epochs

Observations

- Increasing filters improved feature extraction
- Too many layers increased overfitting
- Optimal performance achieved with 2–3 convolution layers
- Adam optimizer provided stable convergence

Impact

- Improved model accuracy
- Better feature extraction
- Reduced overfitting with proper configuration

Conclusion

The CNN model successfully classified images from the CIFAR-10 dataset by learning spatial features through convolution and pooling operations. The experiment demonstrated the effectiveness of CNNs for image classification tasks, especially for complex real-world data, while highlighting the need for careful tuning and computational resources.